

# regression\_forest\_testing

December 14, 2025

## 1 Testing Notebook - Regression Random Forest

```
[1]: from scripts.regression_random_forest import RandomForestRegressor
from sklearn.ensemble import RandomForestRegressor as ↵
    SklearnRandomForestRegressor
from sklearn.model_selection import train_test_split, cross_validate
from sklearn.datasets import make_regression
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

RANDOM_STATE = 42
```

### 1.1 Helper functions

```
[2]: """Generate synthetic regression data using scikit-learn's make_regression_
        function."""
def generate_regression_data(n_samples=100, n_features=1, noise=0.1, ↵
    random_state=RANDOM_STATE):
    X, y = make_regression(n_samples=n_samples, n_features=n_features, ↵
        noise=noise, random_state=random_state)
    print("Generated regression data with {} samples and {} features.".format(n_samples, n_features))
    print("First 5 samples of X:\n", X[:5])
    print("First 5 samples of y:\n", y[:5])
    return X, y

"""Evaluate a regression model using cross-validation and return performance_
        metrics."""
def evaluate_model(model, X, y, cv=5):
    cv_results = cross_validate(
        model,
        X,
        y,
        cv=cv,
        scoring=["r2", "neg_mean_squared_error", "neg_mean_absolute_error"],
```

```

        return_train_score=False,
    )

    return {
        "r2_mean": np.mean(cv_results["test_r2"]),
        "r2_std": np.std(cv_results["test_r2"]),
        "mse_mean": -np.mean(cv_results["test_neg_mean_squared_error"]),
        "mae_mean": -np.mean(cv_results["test_neg_mean_absolute_error"]),
    }
}

```

## 1.2 Definition of testing configurations

```
[3]: TESTING_CONFIGS = [
{
    "description": "Basic regression data with low noise",
    "n_samples": 100,
    "n_features": 1,
    "noise": 0.1
},
{
    "description": "Regression data with higher noise",
    "n_samples": 100,
    "n_features": 1,
    "noise": 10.0
},
{
    "description": "Regression data with multiple features",
    "n_samples": 200,
    "n_features": 5,
    "noise": 5.0
},
{
    "description": "Larger dataset with moderate noise",
    "n_samples": 5000,
    "n_features": 3,
    "noise": 2.0
},
{
    "description": "Large dataset with low noise",
    "n_samples": 10000,
    "n_features": 2,
    "noise": 0.5
},
{
    "description": "Large dataset with high noise",
    "n_samples": 10000,
    "n_features": 2,
}
```

```

    "noise": 20.0
},
{
  "description": "Moderate dataset with many features",
  "n_samples": 1000,
  "n_features": 50,
  "noise": 5.0
}
]
]

HYPERPARAM_CONFIGS = [
{
  "name": "rf_shallow",
  "n_estimators": 100,
  "max_depth": 2,
  "min_samples_leaf": 1,
  "min_samples_split": 2,
  "max_features": "sqrt",
  "bootstrap": True,
},
{
  "name": "rf_medium",
  "n_estimators": 100,
  "max_depth": 4,
  "min_samples_leaf": 1,
  "min_samples_split": 2,
  "max_features": "sqrt",
  "bootstrap": True,
},
{
  "name": "rf_deep",
  "n_estimators": 100,
  "max_depth": 8,
  "min_samples_leaf": 1,
  "min_samples_split": 2,
  "max_features": "sqrt",
  "bootstrap": True,
},
{
  "name": "rf_unrestricted",
  "n_estimators": 100,
  "max_depth": None,
  "min_samples_leaf": 1,
  "min_samples_split": 2,
  "max_features": "sqrt",
  "bootstrap": True,
}
],

```

```

{
    "name": "rf_regularized_leaf",
    "n_estimators": 100,
    "max_depth": None,
    "min_samples_leaf": 5,
    "min_samples_split": 2,
    "max_features": "sqrt",
    "bootstrap": True,
},
{
    "name": "rf_strongly_regularized",
    "n_estimators": 100,
    "max_depth": None,
    "min_samples_leaf": 10,
    "min_samples_split": 10,
    "max_features": "sqrt",
    "bootstrap": True,
},
{
    "name": "rf_bagging_only",
    "n_estimators": 100,
    "max_depth": None,
    "min_samples_leaf": 1,
    "min_samples_split": 2,
    "max_features": None,
    "bootstrap": True,
}
]

```

### 1.3 Runing the tests

For each configuration we will train the scikit-learn's built-in regression random forest model as well as our implementation, we collect useful statistics and compare them.

```
[4]: results = []

for config in TESTING_CONFIGS:
    print("=" * 80)
    print(f"Dataset: {config['description']}")  

    print("=" * 80)

    X, y = generate_regression_data(
        n_samples=config["n_samples"],
        n_features=config["n_features"],
        noise=config["noise"],
    )
```

```

for hp in HYPERPARAM_CONFIGS:
    print("-" * 80)
    print(f"Hyperparams: {hp['name']}")

model_specs = [
    (
        "our",
        RandomForestRegressor(
            n_estimators=hp["n_estimators"],
            max_depth=hp["max_depth"],
            min_samples_leaf=hp["min_samples_leaf"],
            min_samples_split=hp["min_samples_split"],
            max_features=hp["max_features"],
            bootstrap=hp["bootstrap"],
            random_state=RANDOM_STATE
        ),
    ),
    (
        "sklearn",
        SklearnRandomForestRegressor(
            n_estimators=hp["n_estimators"],
            max_depth=hp["max_depth"],
            min_samples_leaf=hp["min_samples_leaf"],
            min_samples_split=hp["min_samples_split"],
            max_features=hp["max_features"],
            bootstrap=hp["bootstrap"],
            random_state=RANDOM_STATE
        ),
    ),
),
]

for model_name, model in model_specs:
    metrics = evaluate_model(model, X, y)

    print(
        f"{model_name:8s} | "
        f"R²: {metrics['r2_mean']:.3f} ± {metrics['r2_std']:.3f} | "
        f"MSE: {metrics['mse_mean']:.3f} | "
        f"MAE: {metrics['mae_mean']:.3f}"
    )

    results.append({
        "dataset": config["description"],
        "hyperparams": hp["name"],
        "model": model_name,
        **metrics,
    })

```

```
=====
Dataset: Basic regression data with low noise
=====
Generated regression data with 100 samples and 1 features.
First 5 samples of X:
[[ 0.93128012]
 [ 0.08704707]
 [-1.05771093]
 [ 0.31424733]
 [-0.47917424]]
First 5 samples of y:
[ 38.9917296   3.4964533  -44.05770173  13.09112657 -19.9786311 ]
-----
Hyperparams: rf_shallow
our      | R2: 0.941 ± 0.029 | MSE: 83.496 | MAE: 6.466
sklearn  | R2: 0.941 ± 0.028 | MSE: 82.845 | MAE: 6.230
-----
Hyperparams: rf_medium
our      | R2: 0.990 ± 0.009 | MSE: 15.910 | MAE: 2.156
sklearn  | R2: 0.990 ± 0.009 | MSE: 15.183 | MAE: 2.099
-----
Hyperparams: rf_deep
our      | R2: 0.993 ± 0.009 | MSE: 11.851 | MAE: 1.334
sklearn  | R2: 0.993 ± 0.009 | MSE: 11.250 | MAE: 1.274
-----
Hyperparams: rf_unrestricted
our      | R2: 0.993 ± 0.009 | MSE: 11.850 | MAE: 1.334
sklearn  | R2: 0.993 ± 0.009 | MSE: 11.250 | MAE: 1.274
-----
Hyperparams: rf_regularized_leaf
our      | R2: 0.985 ± 0.018 | MSE: 25.583 | MAE: 2.354
sklearn  | R2: 0.970 ± 0.025 | MSE: 48.628 | MAE: 3.943
-----
Hyperparams: rf_strongly_regularized
our      | R2: 0.949 ± 0.029 | MSE: 75.613 | MAE: 5.365
sklearn  | R2: 0.915 ± 0.026 | MSE: 122.739 | MAE: 7.441
-----
Hyperparams: rf_bagging_only
our      | R2: 0.993 ± 0.009 | MSE: 11.850 | MAE: 1.334
sklearn  | R2: 0.993 ± 0.009 | MSE: 11.250 | MAE: 1.274
=====
Dataset: Regression data with higher noise
=====
Generated regression data with 100 samples and 1 features.
First 5 samples of X:
[[ 0.93128012]
 [ 0.08704707]
 [-1.05771093]
```

```
[ 0.31424733]  
[-0.47917424]]  
First 5 samples of y:  
[ 50.77992943 -10.06527016 -34.91839191  10.52674299 -17.73837724]
```

---

```
Hyperparams: rf_shallow  
our | R2: 0.877 ± 0.068 | MSE: 192.074 | MAE: 10.672  
sklearn | R2: 0.876 ± 0.070 | MSE: 192.633 | MAE: 10.757
```

---

```
Hyperparams: rf_medium  
our | R2: 0.925 ± 0.030 | MSE: 116.684 | MAE: 8.703  
sklearn | R2: 0.926 ± 0.030 | MSE: 114.503 | MAE: 8.609
```

---

```
Hyperparams: rf_deep  
our | R2: 0.914 ± 0.032 | MSE: 133.687 | MAE: 9.084  
sklearn | R2: 0.915 ± 0.032 | MSE: 132.492 | MAE: 9.097
```

---

```
Hyperparams: rf_unrestricted  
our | R2: 0.914 ± 0.032 | MSE: 134.203 | MAE: 9.105  
sklearn | R2: 0.915 ± 0.032 | MSE: 132.799 | MAE: 9.111
```

---

```
Hyperparams: rf_regularized_leaf  
our | R2: 0.919 ± 0.034 | MSE: 129.923 | MAE: 8.986  
sklearn | R2: 0.908 ± 0.039 | MSE: 150.154 | MAE: 9.487
```

---

```
Hyperparams: rf_strongly_regularized  
our | R2: 0.888 ± 0.052 | MSE: 181.065 | MAE: 10.128  
sklearn | R2: 0.841 ± 0.076 | MSE: 258.953 | MAE: 12.212
```

---

```
Hyperparams: rf_bagging_only  
our | R2: 0.914 ± 0.032 | MSE: 134.203 | MAE: 9.105  
sklearn | R2: 0.915 ± 0.032 | MSE: 132.799 | MAE: 9.111
```

---

```
=====
```

Dataset: Regression data with multiple features

---

```
=====
```

Generated regression data with 200 samples and 5 features.

First 5 samples of X:

```
[[ -0.3853136  0.1990597 -0.60021688  0.46210347  0.06980208]  
[ 0.13074058  1.6324113 -1.43014138 -1.24778318 -0.44004449]  
[-0.77300978  0.22409248  0.0125924 -0.40122047  0.0976761 ]]  
[ -0.57677133 -0.05023811 -0.23894805  0.27045683 -0.90756366]  
[ -0.57581824  0.6141667   0.75750771 -0.2209696 -0.53050115]]
```

First 5 samples of y:

```
[ -23.06825766 -124.34606251     6.99640504 -71.77655611  17.06336023]
```

---

```
Hyperparams: rf_shallow  
our | R2: 0.612 ± 0.034 | MSE: 3556.656 | MAE: 44.420  
sklearn | R2: 0.606 ± 0.044 | MSE: 3609.998 | MAE: 44.611
```

---

Hyperparams: rf\_medium  
our | R<sup>2</sup>: 0.787 ± 0.046 | MSE: 1949.748 | MAE: 32.193  
sklearn | R<sup>2</sup>: 0.786 ± 0.029 | MSE: 1966.775 | MAE: 32.803

---

Hyperparams: rf\_deep  
our | R<sup>2</sup>: 0.876 ± 0.036 | MSE: 1135.633 | MAE: 23.644  
sklearn | R<sup>2</sup>: 0.876 ± 0.037 | MSE: 1143.984 | MAE: 24.194

---

Hyperparams: rf\_unrestricted  
our | R<sup>2</sup>: 0.880 ± 0.034 | MSE: 1103.044 | MAE: 23.330  
sklearn | R<sup>2</sup>: 0.876 ± 0.034 | MSE: 1145.143 | MAE: 23.624

---

Hyperparams: rf\_regularized\_leaf  
our | R<sup>2</sup>: 0.841 ± 0.037 | MSE: 1473.872 | MAE: 26.523  
sklearn | R<sup>2</sup>: 0.813 ± 0.026 | MSE: 1710.348 | MAE: 28.939

---

Hyperparams: rf\_strongly\_regularized  
our | R<sup>2</sup>: 0.793 ± 0.028 | MSE: 1909.516 | MAE: 30.493  
sklearn | R<sup>2</sup>: 0.746 ± 0.023 | MSE: 2340.845 | MAE: 33.863

---

Hyperparams: rf\_bagging\_only  
our | R<sup>2</sup>: 0.914 ± 0.024 | MSE: 782.231 | MAE: 19.692  
sklearn | R<sup>2</sup>: 0.912 ± 0.029 | MSE: 801.749 | MAE: 20.101

---

Dataset: Larger dataset with moderate noise

---

Generated regression data with 5000 samples and 3 features.

First 5 samples of X:

```
[[ 0.67796997 -1.28472777 -0.33102433]
 [ 1.03138053  0.3881858 -0.97027133]
 [-1.21689671  1.36337651 -0.60515624]
 [-0.54429615 -0.50442268 -1.5198928 ]
 [ 0.2074888   0.44567791  0.42350787]]
```

First 5 samples of y:

```
[ -3.90976845 -46.39345674 -98.15873956 -157.1184525    47.56364341]
```

---

Hyperparams: rf\_shallow  
our | R<sup>2</sup>: 0.502 ± 0.014 | MSE: 4625.518 | MAE: 49.928  
sklearn | R<sup>2</sup>: 0.535 ± 0.011 | MSE: 4315.775 | MAE: 47.784

---

Hyperparams: rf\_medium  
our | R<sup>2</sup>: 0.781 ± 0.017 | MSE: 2038.090 | MAE: 32.157  
sklearn | R<sup>2</sup>: 0.794 ± 0.012 | MSE: 1915.485 | MAE: 31.155

---

Hyperparams: rf\_deep  
our | R<sup>2</sup>: 0.950 ± 0.005 | MSE: 465.040 | MAE: 14.148  
sklearn | R<sup>2</sup>: 0.953 ± 0.004 | MSE: 437.215 | MAE: 13.600

---

---

Hyperparams: rf\_unrestricted  
our | R<sup>2</sup>: 0.986 ± 0.003 | MSE: 133.930 | MAE: 6.248  
sklearn | R<sup>2</sup>: 0.986 ± 0.003 | MSE: 132.733 | MAE: 6.309

---

Hyperparams: rf\_regularized\_leaf  
our | R<sup>2</sup>: 0.979 ± 0.005 | MSE: 196.901 | MAE: 7.593  
sklearn | R<sup>2</sup>: 0.974 ± 0.006 | MSE: 239.974 | MAE: 8.378

---

Hyperparams: rf\_strongly\_regularized  
our | R<sup>2</sup>: 0.970 ± 0.006 | MSE: 280.350 | MAE: 9.148  
sklearn | R<sup>2</sup>: 0.962 ± 0.008 | MSE: 349.487 | MAE: 10.378

---

Hyperparams: rf\_bagging\_only  
our | R<sup>2</sup>: 0.997 ± 0.000 | MSE: 23.490 | MAE: 2.804  
sklearn | R<sup>2</sup>: 0.998 ± 0.001 | MSE: 22.924 | MAE: 2.795

---

=====  
Dataset: Large dataset with low noise  
=====

Generated regression data with 10000 samples and 2 features.

First 5 samples of X:

```
[[-0.5691482  1.59040357]  
 [ 0.73487779  0.49097495]  
 [ 0.20069869  1.10623156]  
 [-1.07774393 -0.54427443]  
 [ 0.10507597 -0.88502863]]
```

First 5 samples of y:

```
[ -47.49496686   75.02051971   25.36349671 -108.87452141    5.49344142]
```

---

Hyperparams: rf\_shallow  
our | R<sup>2</sup>: 0.700 ± 0.008 | MSE: 2887.820 | MAE: 39.491  
sklearn | R<sup>2</sup>: 0.717 ± 0.006 | MSE: 2718.175 | MAE: 38.175

---

Hyperparams: rf\_medium  
our | R<sup>2</sup>: 0.911 ± 0.005 | MSE: 853.817 | MAE: 20.397  
sklearn | R<sup>2</sup>: 0.915 ± 0.005 | MSE: 820.737 | MAE: 20.132

---

Hyperparams: rf\_deep  
our | R<sup>2</sup>: 0.992 ± 0.002 | MSE: 78.464 | MAE: 5.565  
sklearn | R<sup>2</sup>: 0.991 ± 0.002 | MSE: 85.436 | MAE: 5.835

---

Hyperparams: rf\_unrestricted  
our | R<sup>2</sup>: 0.999 ± 0.001 | MSE: 13.858 | MAE: 1.367  
sklearn | R<sup>2</sup>: 0.999 ± 0.001 | MSE: 14.386 | MAE: 1.383

---

Hyperparams: rf\_regularized\_leaf  
our | R<sup>2</sup>: 0.998 ± 0.001 | MSE: 22.540 | MAE: 1.731  
sklearn | R<sup>2</sup>: 0.997 ± 0.001 | MSE: 27.004 | MAE: 1.945

---

```
-----  
Hyperparams: rf_strongly_regularized  
our      | R2: 0.997 ± 0.001 | MSE: 33.867 | MAE: 2.218  
sklearn  | R2: 0.995 ± 0.002 | MSE: 47.927 | MAE: 2.818  
-----  
Hyperparams: rf_bagging_only  
our      | R2: 1.000 ± 0.000 | MSE: 1.792 | MAE: 0.712  
sklearn  | R2: 1.000 ± 0.000 | MSE: 1.899 | MAE: 0.711  
=====  
Dataset: Large dataset with high noise  
=====  
Generated regression data with 10000 samples and 2 features.  
First 5 samples of X:  
[[-0.5691482  1.59040357]  
 [ 0.73487779  0.49097495]  
 [ 0.20069869  1.10623156]  
 [-1.07774393 -0.54427443]  
 [ 0.10507597 -0.88502863]]  
First 5 samples of y:  
[-56.9869474   68.99765941   2.79332014 -93.43973055   7.91252206]  
-----  
Hyperparams: rf_shallow  
our      | R2: 0.679 ± 0.010 | MSE: 3220.896 | MAE: 41.475  
sklearn  | R2: 0.696 ± 0.008 | MSE: 3042.469 | MAE: 40.270  
-----  
Hyperparams: rf_medium  
our      | R2: 0.880 ± 0.004 | MSE: 1199.742 | MAE: 25.491  
sklearn  | R2: 0.885 ± 0.005 | MSE: 1151.316 | MAE: 25.076  
-----  
Hyperparams: rf_deep  
our      | R2: 0.953 ± 0.001 | MSE: 471.854 | MAE: 17.063  
sklearn  | R2: 0.953 ± 0.002 | MSE: 466.333 | MAE: 17.008  
-----  
Hyperparams: rf_unrestricted  
our      | R2: 0.953 ± 0.001 | MSE: 470.615 | MAE: 17.266  
sklearn  | R2: 0.953 ± 0.001 | MSE: 470.539 | MAE: 17.255  
-----  
Hyperparams: rf_regularized_leaf  
our      | R2: 0.956 ± 0.001 | MSE: 445.514 | MAE: 16.752  
sklearn  | R2: 0.956 ± 0.001 | MSE: 443.462 | MAE: 16.685  
-----  
Hyperparams: rf_strongly_regularized  
our      | R2: 0.956 ± 0.001 | MSE: 443.745 | MAE: 16.663  
sklearn  | R2: 0.955 ± 0.001 | MSE: 448.253 | MAE: 16.663  
-----  
Hyperparams: rf_bagging_only  
our      | R2: 0.953 ± 0.001 | MSE: 470.577 | MAE: 17.347  
sklearn  | R2: 0.953 ± 0.001 | MSE: 469.645 | MAE: 17.351
```

```
=====
Dataset: Moderate dataset with many features
=====
```

```
Generated regression data with 1000 samples and 50 features.
```

```
First 5 samples of X:
```

```
[[ 0.22210537  0.62386678 -0.16605612  0.63463021  0.75376893  1.65859202
-0.70189811  0.18361217 -2.08195404  0.60286295  0.14837528 -1.13521912
0.9986041  -0.49461035 -0.51024062  0.93668094 -0.399035  -0.38356556
1.1898542  -0.5212666  -1.16138215  1.72223231  0.36539182  1.33601969
0.78801999 -0.83774995  0.56658811  0.92829461 -0.38674766  0.6208172
0.19797213  1.33836855  0.02365751 -2.52871576  1.75988802 -1.207332
-0.82470391 -0.68112696  0.6151632   0.35782542 -0.7864838   0.10673978
0.69987124 -0.03436237 -0.11283199 -0.36389558  0.10516231 -0.53678647
-0.0572436  -0.30034055]
[ 0.83057101 -1.71669373 -0.15501343  1.11435578 -0.82530382  0.60915961
-0.06801749 -0.71940537 -0.05022968 -1.14378475  0.78259513 -0.10643051
1.68403763 -0.58623557 -0.40156696 -0.59052457  0.96258821  0.79276514
-0.16365711  0.99107132 -1.08776862 -0.43419762 -1.11047912 -1.56341347
-0.20934306  0.19477547  1.77209275 -0.18015897  1.54042895  1.66236667
0.17283733  0.34237509  1.7280638  -0.85290707 -1.17369964  0.55164771
-1.65811325  0.11898642  0.20540077 -0.43031452  2.22904372 -1.44808321
-1.14136697 -0.10875609 -0.51455358 -0.14546159  0.740684   0.85154814
-0.72018801 -0.39993873]
[-0.96335321  0.5437933  -1.11313988  1.21235947  0.55206656 -0.91026463
-0.85609014  0.53390779  0.38681169  0.25747809  0.40567447 -1.38173128
-1.34554208  2.22468763  0.15646961 -1.15238917 -0.11994209 -0.0888014
-1.33626712 -2.46829003  0.73874622 -0.74254137  0.21518837  2.30103361
0.6396049  -1.00437615  0.009012   1.53905609  0.68941732 -0.81846411
-0.33594849  1.13719016 -0.84854273 -0.92413914 -0.1971291  -1.59092134
-0.20260583  1.09189461 -1.43450227 -1.83909427  0.63691302  1.81273542
-0.4661267  0.98463866 -0.82158042 -0.80044475 -1.55415405 -0.79595511
0.85645506 -0.13711049]
[ 2.1498434 -0.10467932 -0.43722006 -0.33944537  1.17827468  0.60539814
0.62604166  0.02681872 -0.59611361 -0.325793  -0.26361207  1.89754169
-0.40314991 -0.14655815  1.29817434  0.3420916  -0.23200561  1.83926104
-1.72979741  2.16414981  1.01110931 -0.68284508  0.84036538  1.90220597
2.31782164  0.93558607 -0.64628821 -0.42569971 -0.40269403 -1.27318569
1.76253455  1.41378331 -0.04727162  1.15564512 -1.09488559  0.69360613
-0.94454315  0.02231656  0.04946641 -2.29909535 -0.15431771  0.0542324
0.31923321  0.39123279  0.5049892  -0.00548364 -0.58436466  0.76151403
-0.75676622 -1.26037829]
[-0.5372907  0.06883836 -0.49145363 -2.14575896 -0.98773148 -0.43352192
0.31588178  0.46436161 -0.79141506 -0.20844714 -0.74615901 -0.11587905
0.32387858  0.28785406 -0.58537497  1.2066765  -0.70211482  0.95639433
-0.69194153 -0.20713186  0.54609752  0.32999556 -0.36007058  0.95665943
0.22027397  1.11085552  0.40401589  0.95472592 -0.27240938 -1.51210178
0.50589761 -0.90487118 -0.41419898 -0.42887986 -2.38069159 -1.0825947
-0.5624822  0.92588024 -0.8635109  0.42003858 -1.82769117  0.95275321
```

```

-1.41406661 -0.36089524 1.41201599 1.56501888 1.34935979 -0.75915137
-0.91560975 0.18396914]]
First 5 samples of y:
[137.83775667 -3.71055967 145.69973099 -9.68172968 -83.67934132]
-----
Hyperparams: rf_shallow
our | R2: 0.268 ± 0.021 | MSE: 20293.013 | MAE: 112.459
sklearn | R2: 0.305 ± 0.018 | MSE: 19282.187 | MAE: 109.292
-----
Hyperparams: rf_medium
our | R2: 0.466 ± 0.016 | MSE: 14825.161 | MAE: 95.356
sklearn | R2: 0.487 ± 0.019 | MSE: 14245.598 | MAE: 93.528
-----
Hyperparams: rf_deep
our | R2: 0.644 ± 0.017 | MSE: 9892.831 | MAE: 76.654
sklearn | R2: 0.653 ± 0.023 | MSE: 9652.582 | MAE: 75.588
-----
Hyperparams: rf_unrestricted
our | R2: 0.683 ± 0.018 | MSE: 8813.206 | MAE: 71.746
sklearn | R2: 0.695 ± 0.011 | MSE: 8474.058 | MAE: 70.648
-----
Hyperparams: rf_regularized_leaf
our | R2: 0.667 ± 0.020 | MSE: 9259.127 | MAE: 73.524
sklearn | R2: 0.661 ± 0.022 | MSE: 9424.727 | MAE: 74.527
-----
Hyperparams: rf_strongly_regularized
our | R2: 0.631 ± 0.015 | MSE: 10252.295 | MAE: 77.665
sklearn | R2: 0.613 ± 0.017 | MSE: 10755.422 | MAE: 79.742
-----
Hyperparams: rf_bagging_only
our | R2: 0.881 ± 0.014 | MSE: 3289.210 | MAE: 42.954
sklearn | R2: 0.879 ± 0.016 | MSE: 3351.223 | MAE: 43.271

```

## 1.4 Results

```
[5]: df = pd.DataFrame(results)
paired = (
    df.pivot_table(
        index=["dataset", "hyperparams"],
        columns="model",
        values=["r2_mean", "mse_mean", "mae_mean"]
    )
)

paired.columns = ["_".join(col) for col in paired.columns]
paired = paired.reset_index()
```

```

paired["r2_diff"] = paired["r2_mean_our"] - paired["r2_mean_sklearn"]
paired["mse_diff"] = paired["mse_mean_our"] - paired["mse_mean_sklearn"]
paired["mae_diff"] = paired["mae_mean_our"] - paired["mae_mean_sklearn"]

paired.head()

```

```
[5]:
```

	dataset	hyperparams	mae_mean_our	\
0	Basic regression data with low noise	rf_bagging_only	1.333733	
1	Basic regression data with low noise	rf_deep	1.334308	
2	Basic regression data with low noise	rf_medium	2.156377	
3	Basic regression data with low noise	rf_regularized_leaf	2.354010	
4	Basic regression data with low noise	rf_shallow	6.466155	


  

	mae_mean_sklearn	mse_mean_our	mse_mean_sklearn	r2_mean_our	\
0	1.274271	11.849615	11.249964	0.992930	
1	1.274359	11.850647	11.250346	0.992929	
2	2.098668	15.910313	15.182812	0.989953	
3	3.943339	25.582914	48.628196	0.984780	
4	6.230456	83.495577	82.845285	0.940749	


  

	r2_mean_sklearn	r2_diff	mse_diff	mae_diff
0	0.993262	-0.000331	0.599651	0.059462
1	0.993261	-0.000332	0.600301	0.059949
2	0.990378	-0.000425	0.727501	0.057709
3	0.969691	0.015090	-23.045282	-1.589329
4	0.940898	-0.000150	0.650291	0.235699

```
[6]: plt.figure(figsize=(6, 4))
plt.hist(paired["r2_diff"], bins=20, edgecolor="black")
plt.axvline(0, linestyle="--", color="gray")
plt.xlabel("Δ R2 (our - sklearn)")
plt.ylabel("Count")
plt.title("Distribution of R2 differences across all tests")
plt.tight_layout()
plt.show()
```




```
[7]: plt.figure(figsize=(6, 4))
plt.boxplot(
    paired["r2_diff"],
    vert=False,
    showfliers=True
)
plt.axvline(0, linestyle="--", color="gray")
plt.xlabel("Δ R2 (our - sklearn)")
plt.title("Boxplot of R2 differences")
plt.tight_layout()
plt.show()
```




```
[8]: by_hp = (
    paired.groupby("hyperparams")["r2_diff"]
    .agg(["mean", "std"])
    .reset_index()
)

plt.figure(figsize=(7, 4))
plt.bar(
    by_hp["hyperparams"],
    by_hp["mean"],
    yerr=by_hp["std"],
    capsize=5
)
plt.axhline(0, linestyle="--", color="gray")
plt.ylabel("Mean  $\Delta R^2$ ")
plt.title("Mean  $R^2$  difference per hyperparameter regime")
plt.xticks(rotation=30, ha="right")
plt.tight_layout()
plt.show()
```



```
[9]: worst_by_dataset = (
    paired.groupby("dataset")["r2_diff"]
    .apply(lambda x: x.abs().max())
    .sort_values(ascending=False)
)

plt.figure(figsize=(8, 4))
worst_by_dataset.plot(kind="bar")
plt.ylabel("max |Δ R²|")
plt.title("Worst-case R² difference per dataset")
plt.xticks(rotation=30, ha="right")
plt.tight_layout()
plt.show()
```



```
[10]: fig, axes = plt.subplots(1, 2, figsize=(10, 4))

axes[0].hist(paired["mse_diff"], bins=20, edgecolor="black")
axes[0].axvline(0, linestyle="--", color="gray")
axes[0].set_title("Δ MSE distribution")

axes[1].hist(paired["mae_diff"], bins=20, edgecolor="black")
axes[1].axvline(0, linestyle="--", color="gray")
axes[1].set_title("Δ MAE distribution")

plt.tight_layout()
plt.show()
```

