

# regression\_tree\_testing

December 14, 2025

## 1 Testing Notebook - Regression Tree

TODO: - [ ] create random regressions datasets for testing - [ ] run training and prediction on these datasets and collect statistics - [ ] same as 3 but with the scikit regression tree - [ ] try a bunch of different settings, most likely use grid search with cross validation - [ ] visualize the results

```
[1]: from scripts.regression_tree import RegressionTree
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import train_test_split, cross_validate
from sklearn.datasets import make_regression
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

RANDOM_STATE = 42
```

### 1.1 Helper functions

```
[2]: """Generate synthetic regression data using scikit-learn's make_regression function."""
def generate_regression_data(n_samples=100, n_features=1, noise=0.1, random_state=RANDOM_STATE):
    X, y = make_regression(n_samples=n_samples, n_features=n_features, noise=noise, random_state=random_state)
    print("Generated regression data with {} samples and {} features.".format(n_samples, n_features))
    print("First 5 samples of X:\n", X[:5])
    print("First 5 samples of y:\n", y[:5])
    return X, y

"""Evaluate a regression model using cross-validation and return performance metrics."""
def evaluate_model(model, X, y, cv=5):
    cv_results = cross_validate(
        model,
        X,
```

```

        y,
        cv=cv,
        scoring=["r2", "neg_mean_squared_error", "neg_mean_absolute_error"],
        return_train_score=False,
    )

    return {
        "r2_mean": np.mean(cv_results["test_r2"]),
        "r2_std": np.std(cv_results["test_r2"]),
        "mse_mean": -np.mean(cv_results["test_neg_mean_squared_error"]),
        "mae_mean": -np.mean(cv_results["test_neg_mean_absolute_error"]),
    }
}

```

## 1.2 Definition of testing configurations

```
[3]: TESTING_CONFIGS = [
{
    "description": "Basic regression data with low noise",
    "n_samples": 100,
    "n_features": 1,
    "noise": 0.1
},
{
    "description": "Regression data with higher noise",
    "n_samples": 100,
    "n_features": 1,
    "noise": 10.0
},
{
    "description": "Regression data with multiple features",
    "n_samples": 200,
    "n_features": 5,
    "noise": 5.0
},
{
    "description": "Larger dataset with moderate noise",
    "n_samples": 5000,
    "n_features": 3,
    "noise": 2.0
},
{
    "description": "Large dataset with low noise",
    "n_samples": 10000,
    "n_features": 2,
    "noise": 0.5
},
{
}
```

```

"description": "Large dataset with high noise",
"n_samples": 10000,
"n_features": 2,
"noise": 20.0
},
{
"description": "Moderate dataset with many features",
"n_samples": 1000,
"n_features": 50,
"noise": 5.0
}
]
HYPERPARAM_CONFIGS = [
{
    "name": "shallow",
    "max_depth": 2,
    "min_samples_leaf": 1,
    "min_samples_split": 2,
},
{
    "name": "medium",
    "max_depth": 4,
    "min_samples_leaf": 1,
    "min_samples_split": 2,
},
{
    "name": "deep",
    "max_depth": 8,
    "min_samples_leaf": 1,
    "min_samples_split": 2,
},
{
    "name": "unrestricted",
    "max_depth": None,
    "min_samples_leaf": 1,
    "min_samples_split": 2,
},
{
    "name": "regularized_leaf",
    "max_depth": None,
    "min_samples_leaf": 5,
    "min_samples_split": 2,
},
{
    "name": "strongly_regularized",
    "max_depth": None,

```

```

        "min_samples_leaf": 10,
        "min_samples_split": 10,
    },
]

```

### 1.3 Runing the tests

For each configuration we will train the scikit-learn's built-in regression tree model as well as our implementation, we collect useful statistics and compare them.

```
[4]: results = []

for config in TESTING_CONFIGS:
    print("=" * 80)
    print(f"Dataset: {config['description']}")  

    print("=" * 80)

    X, y = generate_regression_data(
        n_samples=config["n_samples"],
        n_features=config["n_features"],
        noise=config["noise"],
    )

    for hp in HYPERPARAM_CONFIGS:
        print("-" * 80)
        print(f"Hyperparams: {hp['name']}")

        model_specs = [
            (
                "our",
                RegressionTree(
                    max_depth=hp["max_depth"],
                    min_samples_leaf=hp["min_samples_leaf"],
                    min_samples_split=hp["min_samples_split"],
                    random_state=RANDOM_STATE
                ),
            ),
            (
                "sklearn",
                DecisionTreeRegressor(
                    max_depth=hp["max_depth"],
                    min_samples_leaf=hp["min_samples_leaf"],
                    min_samples_split=hp["min_samples_split"],
                    random_state=RANDOM_STATE
                ),
            ),
        ],
    ]

```

```

for model_name, model in model_specs:
    metrics = evaluate_model(model, X, y)

    print(
        f"{model_name:8s} | "
        f"R²: {metrics['r2_mean']:.3f} ± {metrics['r2_std']:.3f} | "
        f"MSE: {metrics['mse_mean']:.3f} | "
        f"MAE: {metrics['mae_mean']:.3f}"
    )

results.append({
    "dataset": config["description"],
    "hyperparams": hp["name"],
    "model": model_name,
    **metrics,
})

```

=====

Dataset: Basic regression data with low noise

=====

Generated regression data with 100 samples and 1 features.

First 5 samples of X:

```

[[ 0.93128012]
 [ 0.08704707]
 [-1.05771093]
 [ 0.31424733]
 [-0.47917424]]

```

First 5 samples of y:

```
[ 38.9917296   3.4964533  -44.05770173  13.09112657 -19.9786311 ]
```

-----

Hyperparams: shallow

```

our      | R²: 0.885 ± 0.029 | MSE: 159.505 | MAE: 10.251
sklearn  | R²: 0.885 ± 0.029 | MSE: 159.505 | MAE: 10.251

```

-----

Hyperparams: medium

```

our      | R²: 0.987 ± 0.006 | MSE: 19.322 | MAE: 3.020
sklearn  | R²: 0.987 ± 0.006 | MSE: 19.322 | MAE: 3.020

```

-----

Hyperparams: deep

```

our      | R²: 0.994 ± 0.008 | MSE: 10.395 | MAE: 1.420
sklearn  | R²: 0.994 ± 0.008 | MSE: 10.395 | MAE: 1.420

```

-----

Hyperparams: unrestricted

```

our      | R²: 0.994 ± 0.008 | MSE: 10.395 | MAE: 1.421
sklearn  | R²: 0.994 ± 0.008 | MSE: 10.395 | MAE: 1.421

```

-----

Hyperparams: regularized\_leaf

```
our      | R2: 0.971 ± 0.015 | MSE: 45.195 | MAE: 4.272
sklearn | R2: 0.971 ± 0.015 | MSE: 45.195 | MAE: 4.272
-----
Hyperparams: strongly_regularized
our      | R2: 0.914 ± 0.032 | MSE: 124.116 | MAE: 8.168
sklearn | R2: 0.914 ± 0.032 | MSE: 124.116 | MAE: 8.168
=====
Dataset: Regression data with higher noise
=====
Generated regression data with 100 samples and 1 features.
First 5 samples of X:
[[ 0.93128012]
 [ 0.08704707]
 [-1.05771093]
 [ 0.31424733]
 [-0.47917424]]
First 5 samples of y:
[ 50.77992943 -10.06527016 -34.91839191  10.52674299 -17.73837724]
-----
Hyperparams: shallow
our      | R2: 0.808 ± 0.073 | MSE: 299.307 | MAE: 13.587
sklearn | R2: 0.808 ± 0.073 | MSE: 299.307 | MAE: 13.587
-----
Hyperparams: medium
our      | R2: 0.918 ± 0.028 | MSE: 127.604 | MAE: 9.026
sklearn | R2: 0.918 ± 0.028 | MSE: 127.604 | MAE: 9.026
-----
Hyperparams: deep
our      | R2: 0.885 ± 0.037 | MSE: 177.317 | MAE: 10.900
sklearn | R2: 0.885 ± 0.037 | MSE: 177.317 | MAE: 10.900
-----
Hyperparams: unrestricted
our      | R2: 0.888 ± 0.035 | MSE: 174.039 | MAE: 10.736
sklearn | R2: 0.888 ± 0.035 | MSE: 174.039 | MAE: 10.736
-----
Hyperparams: regularized_leaf
our      | R2: 0.909 ± 0.027 | MSE: 148.618 | MAE: 9.392
sklearn | R2: 0.909 ± 0.027 | MSE: 148.618 | MAE: 9.392
-----
Hyperparams: strongly_regularized
our      | R2: 0.855 ± 0.047 | MSE: 236.005 | MAE: 11.690
sklearn | R2: 0.855 ± 0.047 | MSE: 236.005 | MAE: 11.690
=====
Dataset: Regression data with multiple features
=====
Generated regression data with 200 samples and 5 features.
First 5 samples of X:
[[-0.3853136  0.1990597 -0.60021688  0.46210347  0.06980208]
```

```

[ 0.13074058  1.6324113 -1.43014138 -1.24778318 -0.44004449]
[-0.77300978  0.22409248  0.0125924 -0.40122047  0.0976761 ]
[-0.57677133 -0.05023811 -0.23894805  0.27045683 -0.90756366]
[-0.57581824  0.6141667   0.75750771 -0.2209696 -0.53050115]]
First 5 samples of y:
[ -23.06825766 -124.34606251     6.99640504   -71.77655611    17.06336023]
-----
Hyperparams: shallow
our      | R2: 0.623 ± 0.079 | MSE: 3325.007 | MAE: 44.673
sklearn  | R2: 0.623 ± 0.079 | MSE: 3325.007 | MAE: 44.673
-----
Hyperparams: medium
our      | R2: 0.766 ± 0.068 | MSE: 2056.622 | MAE: 36.590
sklearn  | R2: 0.759 ± 0.065 | MSE: 2118.765 | MAE: 36.980
-----
Hyperparams: deep
our      | R2: 0.836 ± 0.072 | MSE: 1437.223 | MAE: 29.751
sklearn  | R2: 0.825 ± 0.071 | MSE: 1519.290 | MAE: 29.683
-----
Hyperparams: unrestricted
our      | R2: 0.836 ± 0.066 | MSE: 1435.343 | MAE: 29.745
sklearn  | R2: 0.825 ± 0.063 | MSE: 1545.506 | MAE: 30.016
-----
Hyperparams: regularized_leaf
our      | R2: 0.811 ± 0.043 | MSE: 1709.692 | MAE: 31.868
sklearn  | R2: 0.811 ± 0.043 | MSE: 1709.692 | MAE: 31.868
-----
Hyperparams: strongly_regularized
our      | R2: 0.783 ± 0.014 | MSE: 1974.617 | MAE: 34.900
sklearn  | R2: 0.783 ± 0.014 | MSE: 1974.617 | MAE: 34.900
=====
Dataset: Larger dataset with moderate noise
=====
Generated regression data with 5000 samples and 3 features.
First 5 samples of X:
[[ 0.67796997 -1.28472777 -0.33102433]
 [ 1.03138053  0.3881858 -0.97027133]
 [-1.21689671  1.36337651 -0.60515624]
 [-0.54429615 -0.50442268 -1.5198928 ]
 [ 0.2074888   0.44567791  0.42350787]]
First 5 samples of y:
[ -3.90976845 -46.39345674 -98.15873956 -157.1184525   47.56364341]
-----
Hyperparams: shallow
our      | R2: 0.728 ± 0.016 | MSE: 2520.708 | MAE: 39.683
sklearn  | R2: 0.728 ± 0.016 | MSE: 2520.708 | MAE: 39.683
-----
Hyperparams: medium

```

```
our      | R2: 0.904 ± 0.007 | MSE: 890.915 | MAE: 23.495
sklearn | R2: 0.904 ± 0.007 | MSE: 890.915 | MAE: 23.495
```

---

Hyperparams: deep

```
our      | R2: 0.989 ± 0.001 | MSE: 101.228 | MAE: 7.401
sklearn | R2: 0.989 ± 0.001 | MSE: 98.149 | MAE: 7.354
```

---

Hyperparams: unrestricted

```
our      | R2: 0.994 ± 0.001 | MSE: 55.475 | MAE: 4.884
sklearn | R2: 0.994 ± 0.001 | MSE: 53.924 | MAE: 4.859
```

---

Hyperparams: regularized\_leaf

```
our      | R2: 0.992 ± 0.001 | MSE: 70.037 | MAE: 5.411
sklearn | R2: 0.992 ± 0.001 | MSE: 70.044 | MAE: 5.411
```

---

Hyperparams: strongly\_regularized

```
our      | R2: 0.989 ± 0.002 | MSE: 99.862 | MAE: 6.621
sklearn | R2: 0.989 ± 0.002 | MSE: 99.862 | MAE: 6.621
```

---

Dataset: Large dataset with low noise

---

Generated regression data with 10000 samples and 2 features.

First 5 samples of X:

```
[[-0.5691482  1.59040357]
 [ 0.73487779  0.49097495]
 [ 0.20069869  1.10623156]
 [-1.07774393 -0.54427443]
 [ 0.10507597 -0.88502863]]
```

First 5 samples of y:

```
[ -47.49496686   75.02051971   25.36349671 -108.87452141    5.49344142]
```

---

Hyperparams: shallow

```
our      | R2: 0.877 ± 0.005 | MSE: 1180.012 | MAE: 27.415
sklearn | R2: 0.877 ± 0.005 | MSE: 1180.012 | MAE: 27.415
```

---

Hyperparams: medium

```
our      | R2: 0.987 ± 0.001 | MSE: 128.172 | MAE: 8.694
sklearn | R2: 0.987 ± 0.001 | MSE: 128.172 | MAE: 8.694
```

---

Hyperparams: deep

```
our      | R2: 0.999 ± 0.000 | MSE: 10.967 | MAE: 2.547
sklearn | R2: 0.999 ± 0.000 | MSE: 10.986 | MAE: 2.547
```

---

Hyperparams: unrestricted

```
our      | R2: 1.000 ± 0.000 | MSE: 2.996 | MAE: 1.118
sklearn | R2: 1.000 ± 0.000 | MSE: 3.170 | MAE: 1.129
```

---

Hyperparams: regularized\_leaf

```

our      | R2: 0.999 ± 0.000 | MSE: 5.449 | MAE: 1.394
sklearn | R2: 0.999 ± 0.000 | MSE: 5.448 | MAE: 1.394
-----
Hyperparams: strongly_regularized
our      | R2: 0.999 ± 0.000 | MSE: 10.863 | MAE: 1.810
sklearn | R2: 0.999 ± 0.000 | MSE: 10.863 | MAE: 1.810
=====
Dataset: Large dataset with high noise
=====
Generated regression data with 10000 samples and 2 features.
First 5 samples of X:
[[ -0.5691482  1.59040357]
 [ 0.73487779  0.49097495]
 [ 0.20069869  1.10623156]
 [-1.07774393 -0.54427443]
 [ 0.10507597 -0.88502863]]
First 5 samples of y:
[-56.9869474   68.99765941   2.79332014 -93.43973055   7.91252206]
-----
Hyperparams: shallow
our      | R2: 0.841 ± 0.004 | MSE: 1592.545 | MAE: 31.659
sklearn | R2: 0.841 ± 0.004 | MSE: 1592.545 | MAE: 31.659
-----
Hyperparams: medium
our      | R2: 0.947 ± 0.002 | MSE: 529.733 | MAE: 18.329
sklearn | R2: 0.947 ± 0.002 | MSE: 529.733 | MAE: 18.329
-----
Hyperparams: deep
our      | R2: 0.955 ± 0.001 | MSE: 446.888 | MAE: 16.879
sklearn | R2: 0.955 ± 0.001 | MSE: 447.420 | MAE: 16.882
-----
Hyperparams: unrestricted
our      | R2: 0.920 ± 0.001 | MSE: 801.002 | MAE: 22.730
sklearn | R2: 0.920 ± 0.002 | MSE: 798.026 | MAE: 22.683
-----
Hyperparams: regularized_leaf
our      | R2: 0.945 ± 0.001 | MSE: 552.718 | MAE: 18.864
sklearn | R2: 0.945 ± 0.001 | MSE: 552.655 | MAE: 18.863
-----
Hyperparams: strongly_regularized
our      | R2: 0.951 ± 0.001 | MSE: 491.906 | MAE: 17.711
sklearn | R2: 0.951 ± 0.001 | MSE: 491.906 | MAE: 17.711
=====
Dataset: Moderate dataset with many features
=====
Generated regression data with 1000 samples and 50 features.
First 5 samples of X:
[[ 0.22210537  0.62386678 -0.16605612  0.63463021  0.75376893  1.65859202

```

```

-0.70189811  0.18361217 -2.08195404  0.60286295  0.14837528 -1.13521912
 0.9986041   -0.49461035 -0.51024062  0.93668094 -0.399035   -0.38356556
 1.1898542   -0.5212666  -1.16138215  1.72223231  0.36539182  1.33601969
 0.78801999 -0.83774995  0.56658811  0.92829461 -0.38674766  0.6208172
 0.19797213  1.33836855  0.02365751 -2.52871576  1.75988802 -1.207332
-0.82470391 -0.68112696  0.6151632   0.35782542 -0.7864838   0.10673978
 0.69987124 -0.03436237 -0.11283199 -0.36389558  0.10516231 -0.53678647
-0.0572436  -0.30034055]
[ 0.83057101 -1.71669373 -0.15501343  1.11435578 -0.82530382  0.60915961
-0.06801749 -0.71940537 -0.05022968 -1.14378475  0.78259513 -0.10643051
 1.68403763 -0.58623557 -0.40156696 -0.59052457  0.96258821  0.79276514
-0.16365711  0.99107132 -1.08776862 -0.43419762 -1.11047912 -1.56341347
-0.20934306  0.19477547  1.77209275 -0.18015897  1.54042895  1.66236667
 0.17283733  0.34237509  1.7280638  -0.85290707 -1.17369964  0.55164771
-1.65811325  0.11898642  0.20540077 -0.43031452  2.22904372 -1.44808321
-1.14136697 -0.10875609 -0.51455358 -0.14546159  0.740684   0.85154814
-0.72018801 -0.39993873]
[-0.96335321  0.5437933  -1.11313988  1.21235947  0.55206656 -0.91026463
-0.85609014  0.53390779  0.38681169  0.25747809  0.40567447 -1.38173128
-1.34554208  2.22468763  0.15646961 -1.15238917 -0.11994209 -0.0888014
-1.33626712 -2.46829003  0.73874622 -0.74254137  0.21518837  2.30103361
 0.6396049   -1.00437615  0.009012   1.53905609  0.68941732 -0.81846411
-0.33594849  1.13719016 -0.84854273 -0.92413914 -0.1971291  -1.59092134
-0.20260583  1.09189461 -1.43450227 -1.83909427  0.63691302  1.81273542
-0.4661267   0.98463866 -0.82158042 -0.80044475 -1.55415405 -0.79595511
 0.85645506 -0.13711049]
[ 2.1498434  -0.10467932 -0.43722006 -0.33944537  1.17827468  0.60539814
 0.62604166  0.02681872 -0.59611361 -0.325793  -0.26361207  1.89754169
-0.40314991 -0.14655815  1.29817434  0.3420916  -0.23200561  1.83926104
-1.72979741  2.16414981  1.01110931 -0.68284508  0.84036538  1.90220597
 2.31782164  0.93558607 -0.64628821 -0.42569971 -0.40269403 -1.27318569
 1.76253455  1.41378331 -0.04727162  1.15564512 -1.09488559  0.69360613
-0.94454315  0.02231656  0.04946641 -2.29909535 -0.15431771  0.0542324
 0.31923321  0.39123279  0.5049892  -0.00548364 -0.58436466  0.76151403
-0.75676622 -1.26037829]
[-0.5372907   0.06883836 -0.49145363 -2.14575896 -0.98773148 -0.43352192
 0.31588178  0.46436161 -0.79141506 -0.20844714 -0.74615901 -0.11587905
 0.32387858  0.28785406 -0.58537497  1.2066765  -0.70211482  0.95639433
-0.69194153 -0.20713186  0.54609752  0.32999556 -0.36007058  0.95665943
 0.22027397  1.11085552  0.40401589  0.95472592 -0.27240938 -1.51210178
 0.50589761 -0.90487118 -0.41419898 -0.42887986 -2.38069159 -1.0825947
-0.5624822   0.92588024 -0.8635109   0.42003858 -1.82769117  0.95275321
-1.41406661 -0.36089524  1.41201599  1.56501888  1.34935979 -0.75915137
-0.91560975  0.18396914]]

```

First 5 samples of y:

```
[137.83775667 -3.71055967 145.69973099 -9.68172968 -83.67934132]
```

---

Hyperparams: shallow

```
our      | R2: 0.380 ± 0.039 | MSE: 17106.506 | MAE: 104.267
sklearn | R2: 0.380 ± 0.039 | MSE: 17106.506 | MAE: 104.267
```

---

Hyperparams: medium

```
our      | R2: 0.670 ± 0.027 | MSE: 9083.463 | MAE: 76.403
sklearn | R2: 0.668 ± 0.028 | MSE: 9145.507 | MAE: 76.560
```

---

Hyperparams: deep

```
our      | R2: 0.727 ± 0.026 | MSE: 7523.640 | MAE: 67.167
sklearn | R2: 0.727 ± 0.037 | MSE: 7488.662 | MAE: 66.500
```

---

Hyperparams: unrestricted

```
our      | R2: 0.705 ± 0.032 | MSE: 8120.275 | MAE: 69.466
sklearn | R2: 0.701 ± 0.039 | MSE: 8223.028 | MAE: 69.967
```

---

Hyperparams: regularized\_leaf

```
our      | R2: 0.756 ± 0.026 | MSE: 6707.463 | MAE: 64.719
sklearn | R2: 0.756 ± 0.027 | MSE: 6696.210 | MAE: 64.661
```

---

Hyperparams: strongly\_regularized

```
our      | R2: 0.758 ± 0.028 | MSE: 6653.061 | MAE: 64.376
sklearn | R2: 0.758 ± 0.028 | MSE: 6653.061 | MAE: 64.376
```

## 1.4 Results

```
[5]: df = pd.DataFrame(results)
paired = (
    df.pivot_table(
        index=["dataset", "hyperparams"],
        columns="model",
        values=["r2_mean", "mse_mean", "mae_mean"]
    )
)

paired.columns = ["_".join(col) for col in paired.columns]
paired = paired.reset_index()

paired["r2_diff"] = paired["r2_mean_our"] - paired["r2_mean_sklearn"]
paired["mse_diff"] = paired["mse_mean_our"] - paired["mse_mean_sklearn"]
paired["mae_diff"] = paired["mae_mean_our"] - paired["mae_mean_sklearn"]

paired.head()
```

```
[5]:          dataset      hyperparams  mae_mean_our \
0  Basic regression data with low noise      deep      1.420082
1  Basic regression data with low noise    medium      3.019620
2  Basic regression data with low noise  regularized_leaf      4.271927
```

```

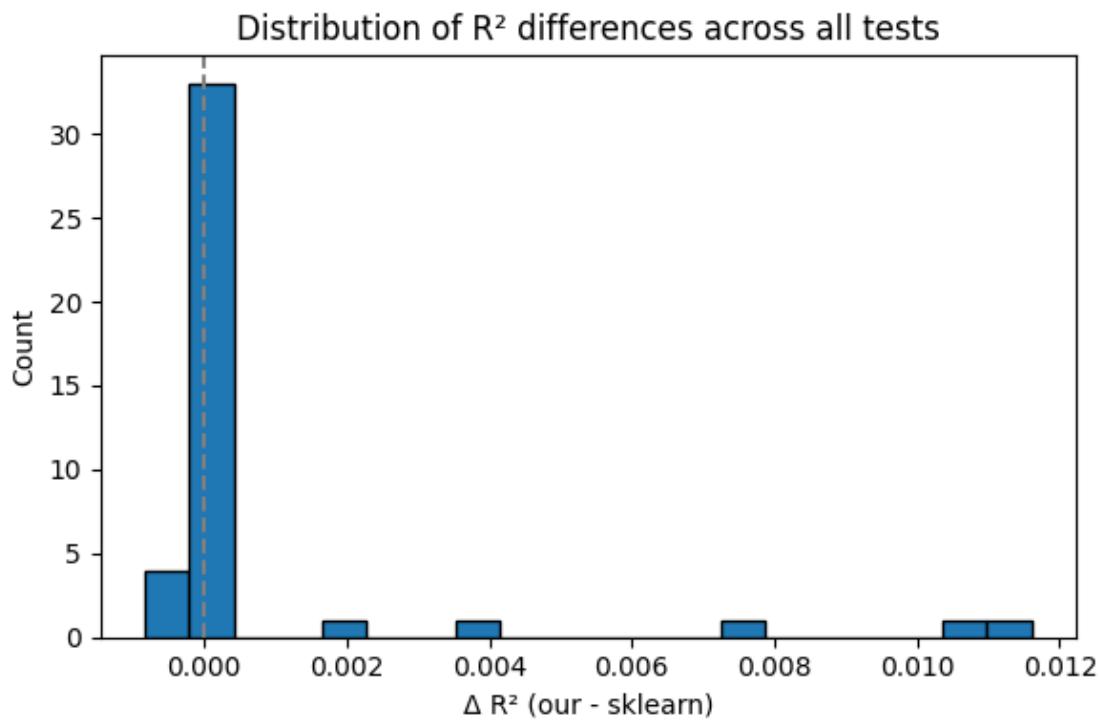
3 Basic regression data with low noise           shallow      10.251087
4 Basic regression data with low noise  strongly_regularized   8.168215

mae_mean_sklearn  mse_mean_our  mse_mean_sklearn  r2_mean_our  \
0            1.420082    10.395394     10.395394    0.993668
1            3.019620    19.322168     19.322168    0.986635
2            4.271927    45.194549     45.194549    0.971219
3           10.251087   159.505327    159.505327    0.884677
4            8.168215   124.116470    124.116470    0.913597

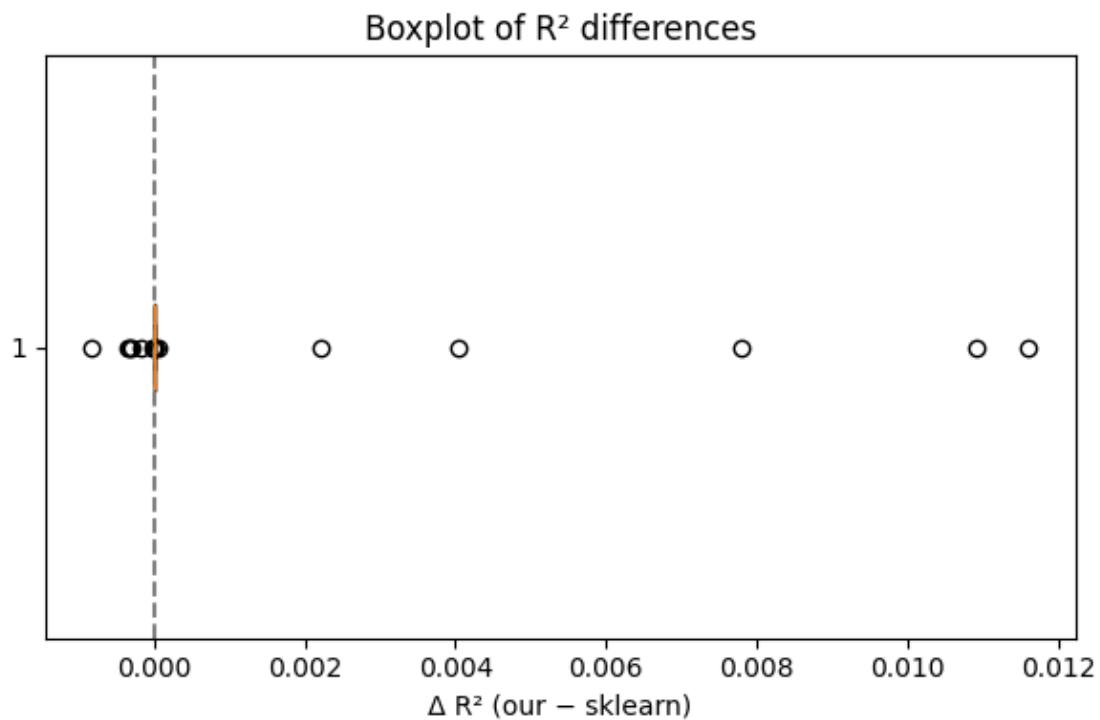
r2_mean_sklearn  r2_diff      mse_diff       mae_diff
0            0.993668      0.0  0.000000e+00  0.000000e+00
1            0.986635      0.0  0.000000e+00  0.000000e+00
2            0.971219      0.0  2.131628e-14  1.776357e-15
3            0.884677      0.0 -5.684342e-14  1.776357e-15
4            0.913597      0.0  0.000000e+00  0.000000e+00

```

```
[6]: plt.figure(figsize=(6, 4))
plt.hist(paired["r2_diff"], bins=20, edgecolor="black")
plt.axvline(0, linestyle="--", color="gray")
plt.xlabel("Δ R2 (our - sklearn)")
plt.ylabel("Count")
plt.title("Distribution of R2 differences across all tests")
plt.tight_layout()
plt.show()
```



```
[7]: plt.figure(figsize=(6, 4))
plt.boxplot(
    paired["r2_diff"],
    vert=False,
    showfliers=True
)
plt.axvline(0, linestyle="--", color="gray")
plt.xlabel("Δ R2 (our - sklearn)")
plt.title("Boxplot of R2 differences")
plt.tight_layout()
plt.show()
```



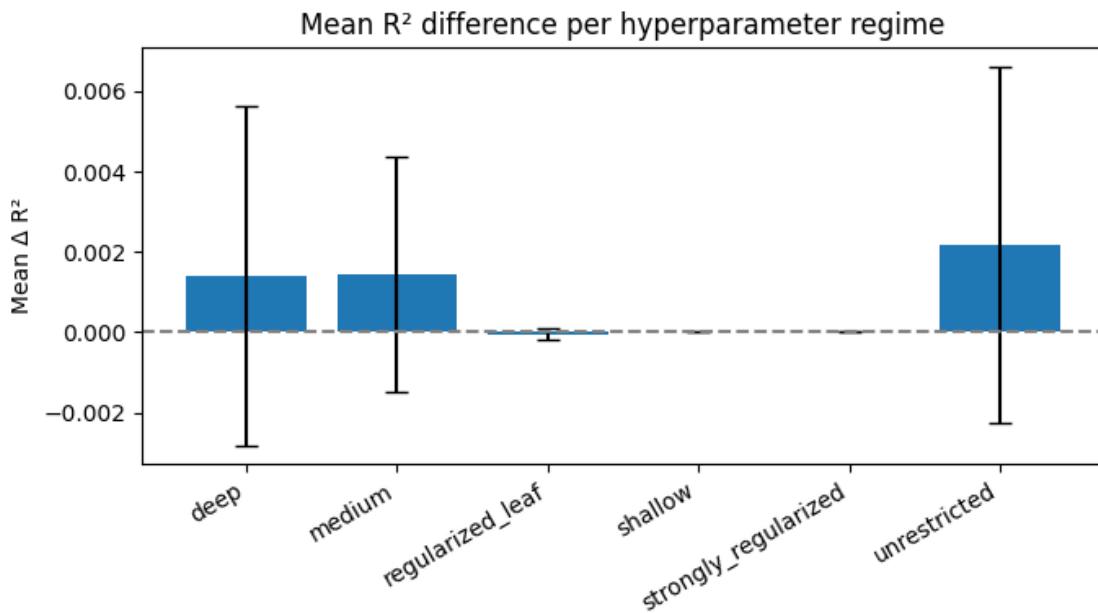
```
[8]: by_hp = (
    paired.groupby("hyperparams")["r2_diff"]
    .agg(["mean", "std"])
    .reset_index()
)

plt.figure(figsize=(7, 4))
plt.bar(
```

```

        by_hp["hyperparams"],
        by_hp["mean"],
        yerr=by_hp["std"],
        capsize=5
    )
plt.axhline(0, linestyle="--", color="gray")
plt.ylabel("Mean  $\Delta R^2$ ")
plt.title("Mean  $R^2$  difference per hyperparameter regime")
plt.xticks(rotation=30, ha="right")
plt.tight_layout()
plt.show()

```

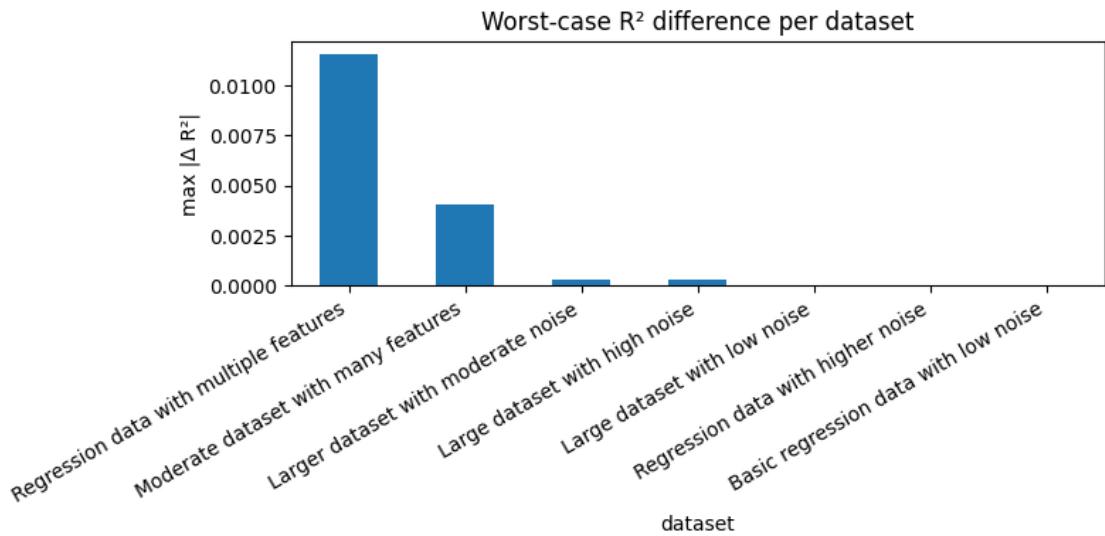


```

[9]: worst_by_dataset = (
    paired.groupby("dataset")["r2_diff"]
    .apply(lambda x: x.abs().max())
    .sort_values(ascending=False)
)

plt.figure(figsize=(8, 4))
worst_by_dataset.plot(kind="bar")
plt.ylabel("max  $|\Delta R^2|$ ")
plt.title("Worst-case  $R^2$  difference per dataset")
plt.xticks(rotation=30, ha="right")
plt.tight_layout()
plt.show()

```



```
[10]: fig, axes = plt.subplots(1, 2, figsize=(10, 4))

axes[0].hist(paired["mse_diff"], bins=20, edgecolor="black")
axes[0].axvline(0, linestyle="--", color="gray")
axes[0].set_title("Δ MSE distribution")

axes[1].hist(paired["mae_diff"], bins=20, edgecolor="black")
axes[1].axvline(0, linestyle="--", color="gray")
axes[1].set_title("Δ MAE distribution")

plt.tight_layout()
plt.show()
```

